


The Bare Metal



09.10.2021

<https://pixabay.com/de/photos/server-rack-server-elektronik-kabel-441494/>

Motivation

- Why even study the „bare metal“ of a computer?
- Your answers here :) ...
- Some observations:
 - CPU: dumb but super fast (ns per instruction)
 - Storage: dumb but super large (ms for disks, μ s SSD, TB)
 - Periphery: very dumb (think mouse), slow (Mbits/s, 8μ s/byte)
 - Network: complicated, slowest but can't do without (Mbit/s, but long latency)
- How to do something useful with such a machine???

How to make a computer work for us

- We can program something to calculate on the CPU.
- For this, we might need data or input from
 - Storage
 - Periphery
 - Network
- The program will often need to WAIT AGES for the data to arrive.
- Meanwhile, we heat the house with the CPU-fan...
- Can we do better? How?

How to keep our programs running

- Suppose we wrote a nice program as indicated on the prev. slide.
- And then, new hardware is available: faster, not so dumb
- But would our program run right out of the box?
- Why, why not?

The role of the OS

- The OS acts as an abstraction and resource manager of the HW.
- It manages
 - Which program may continue running on the CPU for how long
 - Fetching of data from various peripherals (RAM, disk, network)
- It presents the hardware as a consistent „computer“ to programs.
- Therefore,
 - **The OS is the ONLY program directly interacting with the HW!**
 - And it needs support from the HW to do it well!

Content

- What happens on a double click???
- A 10000ft view of a PC
- The x86 Architecture
- Assembler capabilities
- From C to hardware – a long journey (from Java, still longer :))

What happens on a double click?

- The program starts to run...
- What IS an Executable?
 - Human readable code?
 - Machine readable code!
- Where does it run from?
 - Persistent storage → Main memory (RAM)
 - Instructions execute on CPU
- What does a program need to run?
 - Memory, CPU time, peripherals

A 10000ft view of a PC

- CPU
 - Computing power
 - Hardware control
- MMU
 - Memory access
- Internal Peripherals Bus (PCIe)
 - Graphics, Extensions
- External Peripherals Bus
 - SATA, USB



<https://pixabay.com/de/photos/motherboard-waschbecken-ventilator-197608/>

CPU: The x86 Architecture

- Most common PC CPU Architecture
- A couple of onboard memory slots (~40) - **registers**
 - Bus width (64bit)
 - Fastest for calculations
- A built-in hardware based programming language
 - Instruction set (humans write this in assembler; **opcodes; CISC**)
- Provides a stack based programming model
- Provides modes which disables certain instructions (**user mode, kernel mode**)


Oh my...

<code>mov ax, 1</code>	Write 1 to register ax	<code>ax = (byte)1</code>
<code>cwd</code>	Copy ax to dx and make it signed	<code>dx = (int)ax</code>
<code>:Label1</code>	Place a mark to jump to	<code>:Label 1</code>
<code>add dx, ax</code>	Add registers ax and dx, store in dx	<code>dx = dx + ax</code>
<code>inc ax</code>	Increment ax by 1	<code>ax = ax + 1</code>
<code>cmp ax, 10</code>	Compare ax to 10. Sets status flag	<code>sf = ax == 10 ? 1 : 0</code>
<code>jbe Label1</code>	If status flag == 1, jump	<code>if(sf) goto :Label1</code>


How long does this programm need to count to 10 on a 1 GHz CPU?

Interacting with hardware

main:				
sbi DDRB, 5	Set PB5 as output			
blink:	Label			
sbi PINB, 5	Toggle PINB			
ldi r25, hi8(1000)	Load hi-byte into r25			
ldi r24, lo8(1000)	Load lo-byte into r24			
call delay_ms	Jump to delay label			
jmp blink	Jump to blink label			
delay_ms:				
	Delay about (r25:r24)*ms. Clobbers r30, and r31.			
	One millisecond is about 16000 cycles at 16MHz.			
	The inner loop takes 4 cycles, so we repeat it 4000 times			
ldi r31, hi8(4000)	Load hi-byte into r31			
ldi r30, lo8(4000)	Load lo-byte into r30			
1				
sbiw r30, 1	Subtracts 1 from (r31, r30). Sets Zero-flag if result is zero			
brne b1	Branch to b1 if Zero-flag set			
sbiw r24, 1	Subtracts 1 from (r25, r24). Sets Zero-flag if result is zero			
brne delay_ms	Branch to delay_ms if Zero-flag set.			
ret	Return to call			



Alternating LED Blinker: HW



Alternating LED Blinker: SW

```
10  main:
11      sbi    DDRB, 5      ; Set PIN 13 as output
12  → sbi    DDRB, 4      ; Set PIN 12 as output
13  → sbi    PINB, 4      ; Toggle PIN 12
14  blink:
15      sbi    PINB, 5      ; Toggle PIN 13
16  → sbi    PINB, 4      ; Toggle PIN 12
```


From C to hardware

```
#include <stdio.h>
#include <unistd.h>

int main(int argc, char **argv)
{
    char* hello = "Hello, World!";
    write(STDIN_FILENO, hello, 13);
    return 0;
}
```

```
0000000000400507 <main>:
400507: 55                push    %rbp
400508: 48 89 e5          mov     %rsp,%rbp
40050b: 48 83 ec 20       sub     $0x20,%rsp
40050f: 89 7d ec          mov     %edi,-0x14(%rbp)
400512: 48 89 75 e0       mov     %rsi,-0x20(%rbp)
400516: 48 c7 45 f8 c4 05 40 movq    $0x4005c4,-0x8(%rbp)
40051d: 00
40051e: 48 8b 45 f8       mov     -0x8(%rbp),%rax
400522: ba 0d 00 00 00    mov     $0xd,%edx
400527: 48 89 c6          mov     %rax,%rsi
40052a: bf 00 00 00 00    mov     $0x0,%edi
40052f: e8 ec fe ff ff   callq   400420 <write@plt>
400534: b8 00 00 00 00    mov     $0x0,%eax
400539:
40053a:
40053b:
```

```
libc6.txt x  untitled x
ed149: 0f 1f 80 00 00 00 00 nopl    0x0(%rax)
000000000000ed150 <__write>:
ed150: 8b 05 4a d2 2c 00 mov     0x2cd24a(%rip),%eax
ed156: 48 63 ff          movslq  %edi,%rdi
ed159: 85 c0            test    %eax,%eax
ed15b: 75 13            jne     ed170 <__write+0x20>
ed15d: b8 01 00 00 00    mov     $0x1,%eax
ed162: 0f 05            syscall
ed164: 48 3d 00 f0 ff ff cmp     $0xffffffffffffffff000,%rax
ed16a: 77 54            ja      ed1c0 <__write+0x70>
```



Key points

- The **OS** is the **only** program interacting **directly** with **CPU/HW..**
- It provides **System Calls** for programs to access HW.
 - Representation of different HW with same System Calls.
- Normal programs are **NOT ALLOWED** to interact directly with **HW**
 - CPU in **USER MODE**
- Only the OS interacts with HW
 - CPU in **KERNEL MODE**
- Happy hacking :)

The Bare Metal – Your notes



- Hardware eines modernen PC
-
- Rolle des Betriebssystems
-
- CPU: Was kann diese Komponente? Spezielle Modi?
-
- Assembler-Programmierung: Eigenschaften, Struktur
-